Return on Investment and Compute Infrastructure when value of result is not quantified

Ben Payne

20160506

Rapid Prototyping platforms

- Your desktop - Microsoft Excel, etc
- AWS
- Google Compute Engine
- Jupyter notebooks

Triage process design

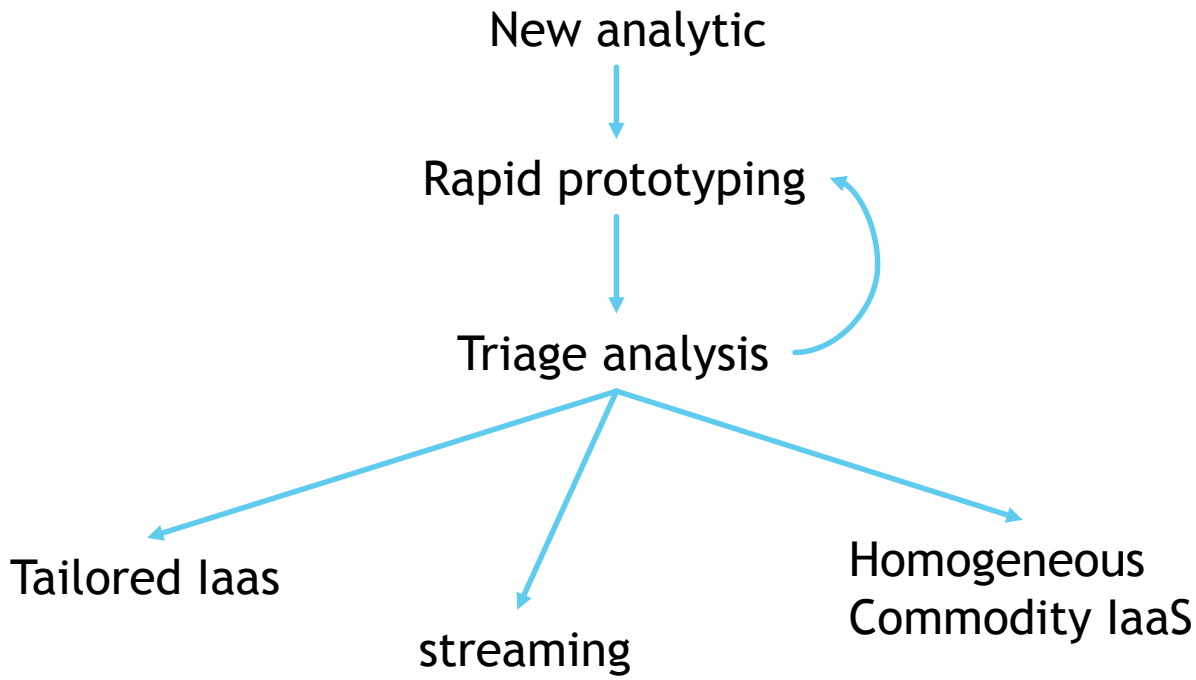

How to evaluate merit?

Two approaches

- Single analytic, compare RoI for multiple platforms over lifespan (ie 5 years)
- Many analytics, compare RoI for a single investment period (ie fiscal year)

Comparison of platforms: “pay-as-you-go” vs “commodity” vs “tailored”

- AWS: assume zero capital cost and zero operations and maintenance cost. Fiscal cost is in number of hours used; I assume a rate of $100/hour. Time-to-market is the advantage; assume 1 day for acquisition and creating the analytic. Time-to-solution is assumed to be 100 minutes

- self-hosted commodity platform: capital cost is $100,000 (one-time investment), and the operations and maintenance is $10,000 per year (an on-going cost). Time-to-market is 5 days; time-to-solution is 60 minutes.

- tailored architecture is costly — $1,000,000 for capital, and $100,000 per year for O&M. Advantage is a tiny time-to-solution, 1 minute, but the time-to-market is 1 year. Time-to-market includes NRE and acquisition.

Total cost of ownership

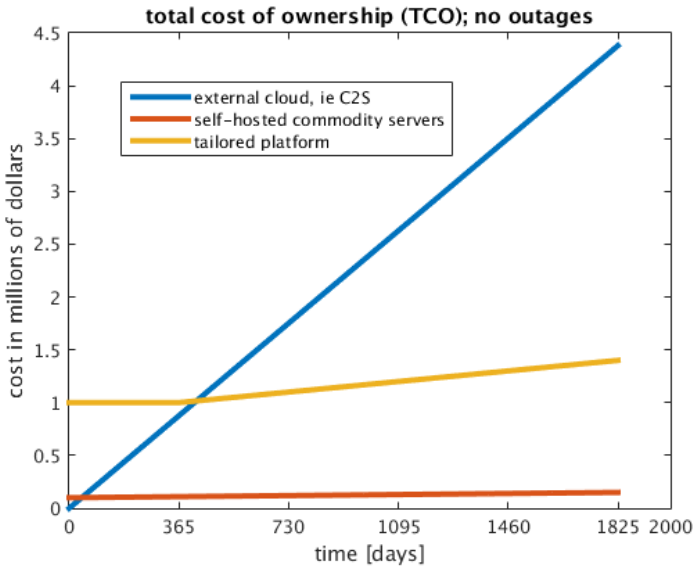

(Lower is better)

(Assume a single analytic)

Observations:

- No capital investment means initial cost is lowest
- Tailored architecture costs more than commodity
- Paying by the hour gets expensive

Total cost of ownership

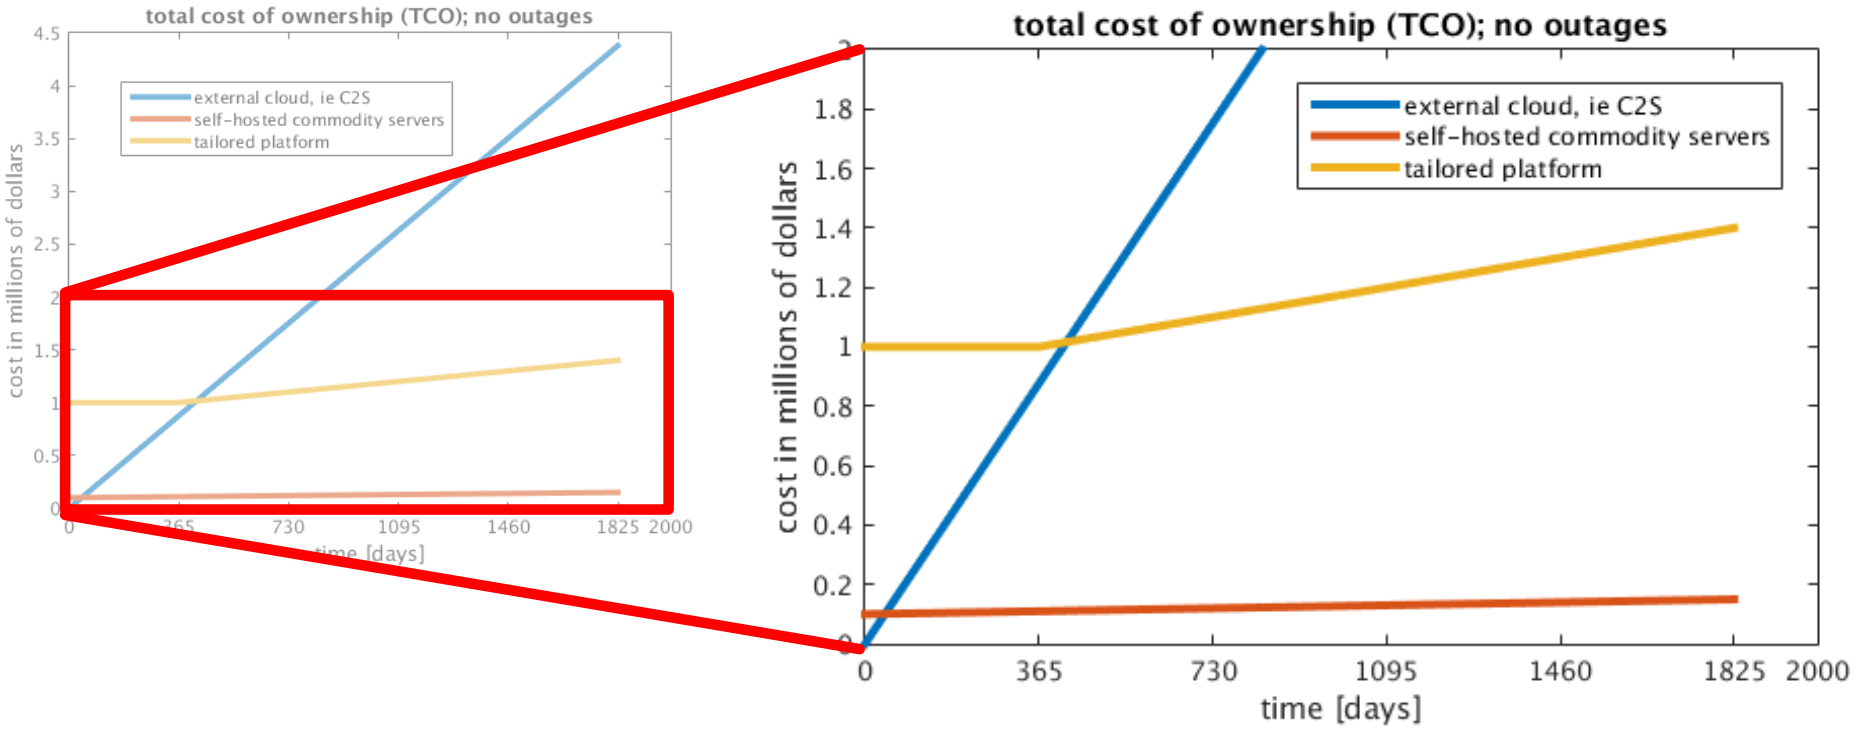

For the current parameters, pay-by-the-hour is not as cost effective;
you do get access to compute faster

Time-to-market and time-to-solution

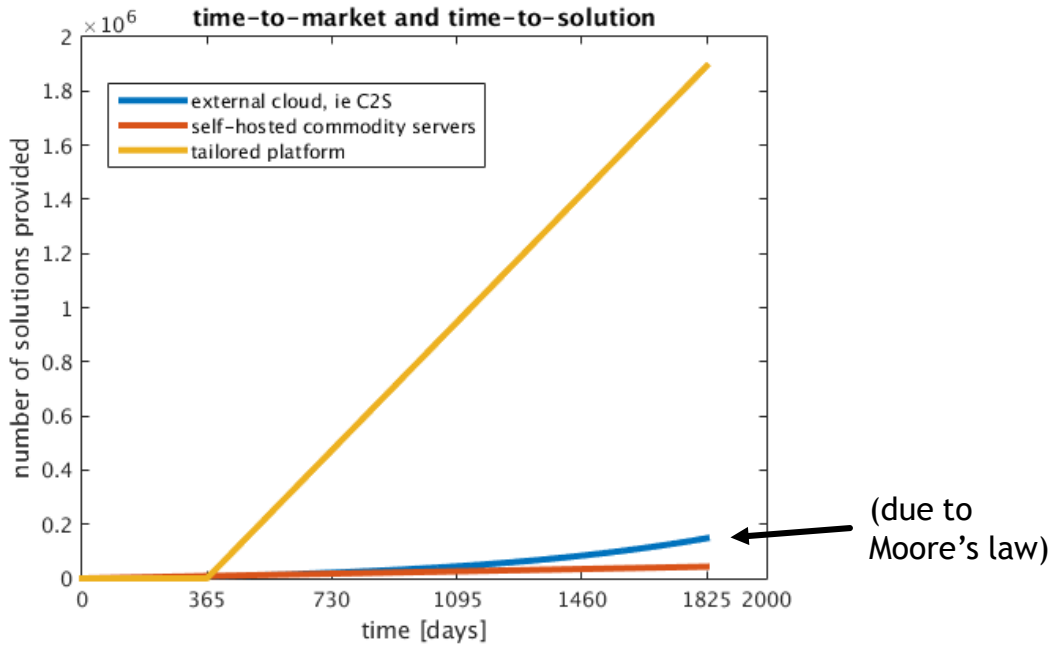

(due to Moore's law)

- The tailored architecture provides more solutions over lifespan

(Higher is better)

System availability:

- External cloud: 99.9%
- Self-hosted cloud: 99%
- Tailored architecture: 90%

Time-to-market and time-to-solution

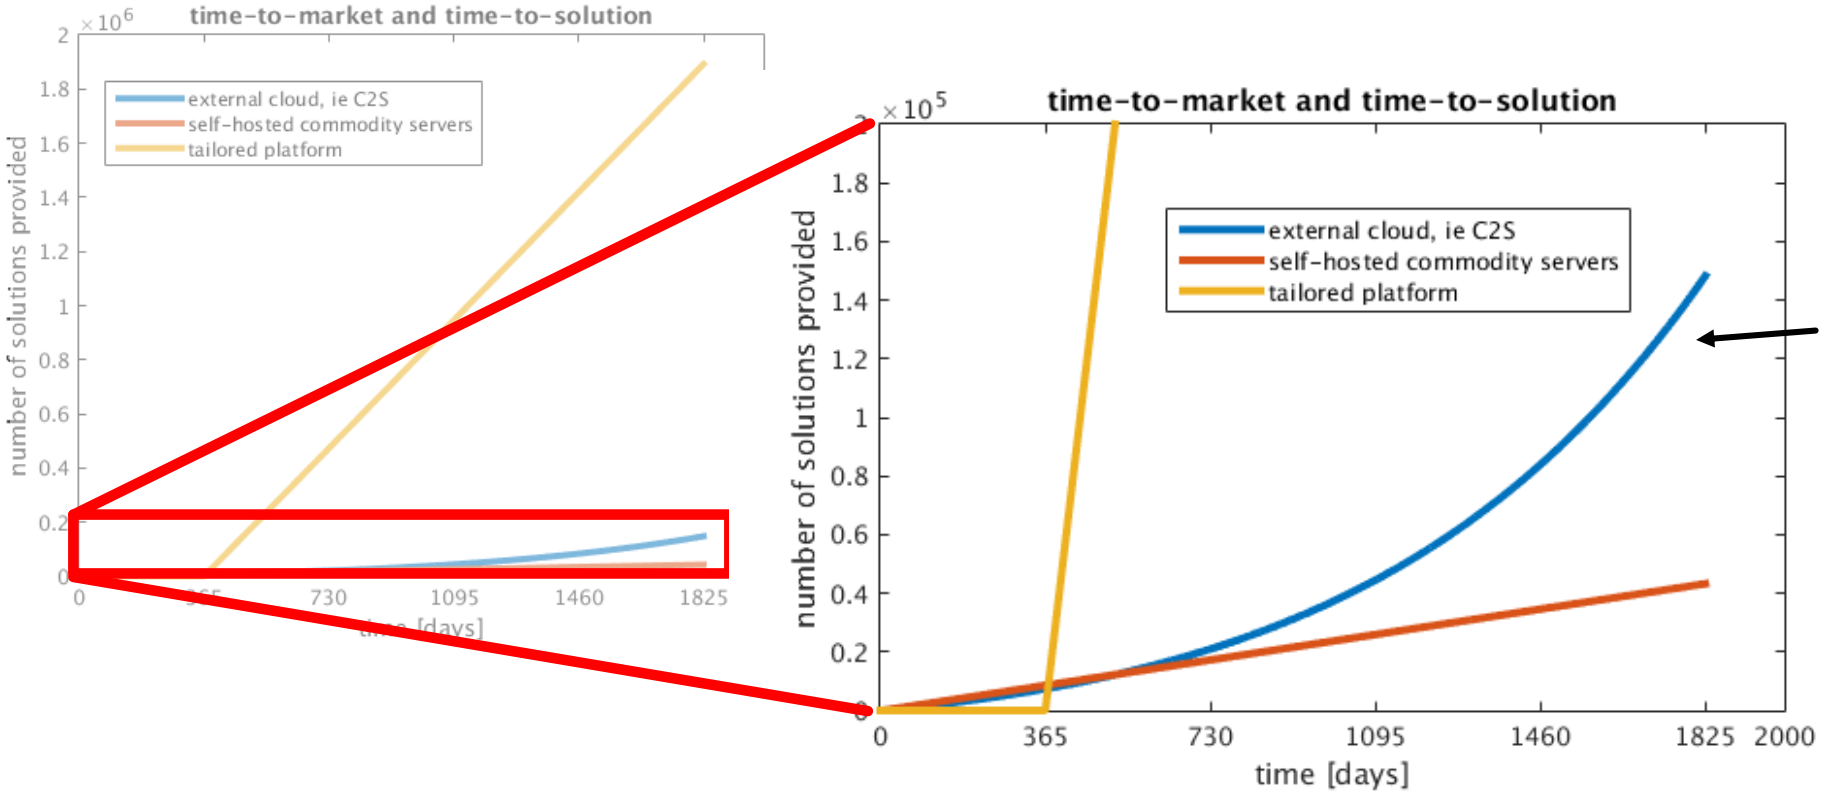

(due to Moore's law)

For the current parameters, ratio of time spent in NRE and acquisition to time-to-solution significantly impacts throughput

Bang-for-your-buck

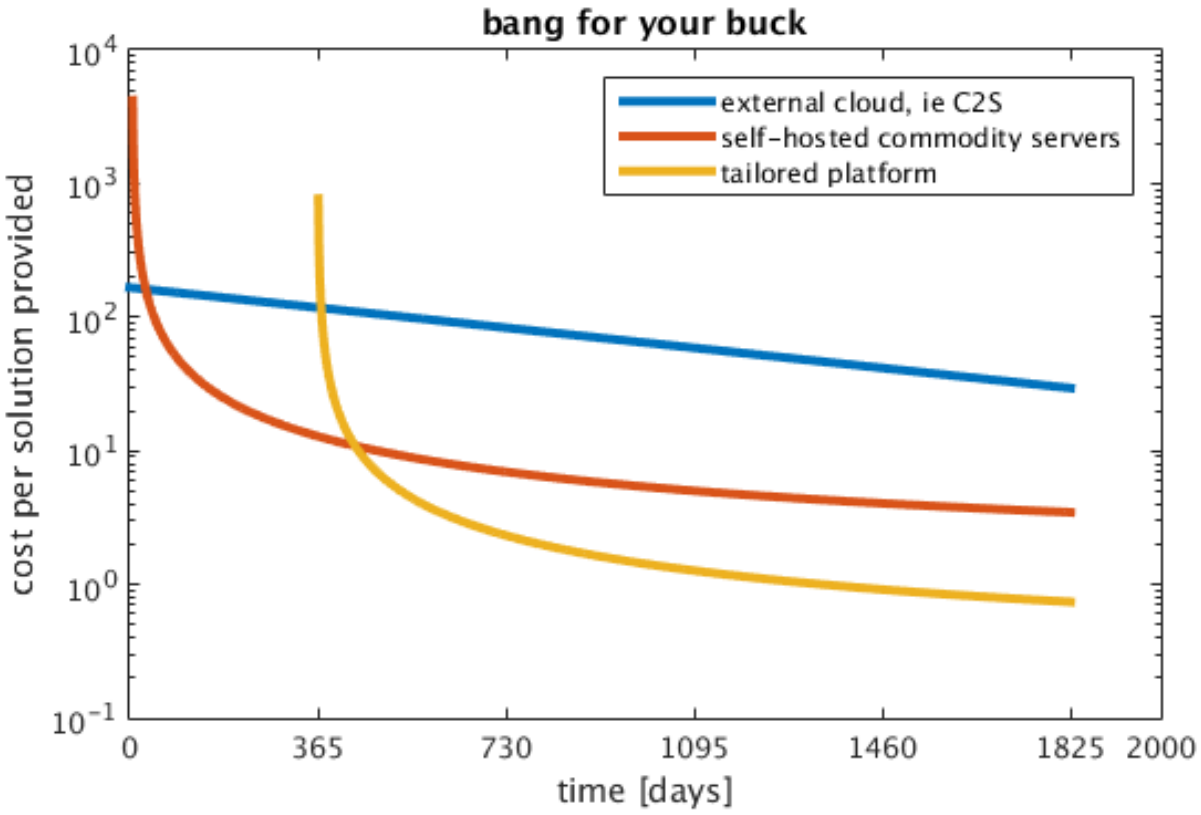

(Lower is better)

Over the lifespan of the system, tailored IaaS is better ROI

Bang-for-your-buck

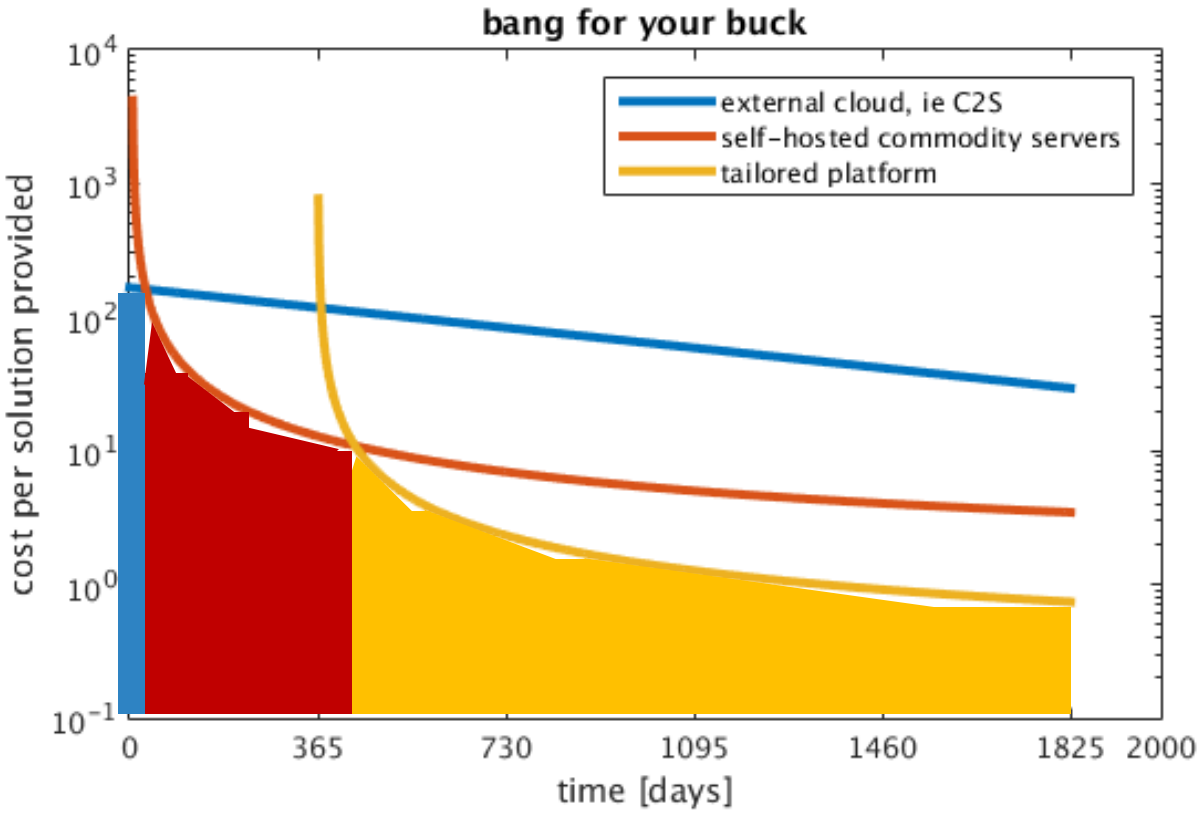

(Lower is better)

"best" depends on durability of analytic

Multi-tenancy on existing infrastructure: how to invest in next iteration?

- Suppose existing platform environment is homogenous
- Suppose you will get money to spend on compute for next fiscal period

- **Should that money be invested on more of the same,**
  or spend the money on a novel architecture?

- Assumption: known tasks in the time period
- Assumption: known amount of work in the period

Input Parameters

- number of existing compute units for architecture A: 5

- capital cost per compute unit for architecture A: $20
- capital cost per compute unit for architecture B: $50
- O&M cost per compute unit for architecture A: $5
- O&M cost per compute unit for architecture B: $8
- Money to invest in next fiscal period: $100

- For each analytic, CPU hours for architecture A: {40, 100, 10, 5, 2}
- For each analytic, CPU hours for architecture B: {30, 10, 8, 4, 1}

Result of analysis

current_tts =

31.4000

tts_homogeneous =

19.6250

min_tts_heterogeneous =

21.4000

Best analytic distribution, platform A =

40 0 10 5 2

best analytic distribution, platform B =

0 10 0 0 0

Result: maintaining homogeneous architecture yields better throughput

current_tts =

31.4000

tts_homogeneous =

19.6250

min_tts_heterogeneous =

21.4000

Best analytic distribution, platform A =

{40 0 10 5 2}

best analytic distribution, platform B =

{0 10 0 0 0}

- For these parameters, change of architecture doesn't outweigh the cost of the acquisition + O&M
- Even though architecture B is always faster for every analytic, the overall mission throughput is lower for the amount of money spent. Therefore, **stick with homogeneous architecture**

Result: different initial conditions yield opposite outcome

capital cost per compute unit for architecture A: $20

current_tts =

31.4000

tts_homogeneous =

19.6250

min_tts_heterogeneous =

21.4000

Best analytic distribution, platform A =

{40 0 10 5 2}

best analytic distribution, platform B =

{0 10 0 0 0}

*capital cost per compute unit for architecture A: **$30***

current_tts =

31.4000

tts_homogeneous =

22.4286

min_tts_heterogeneous =

21.4000

Best analytic distribution, platform A =

{40 0 10 5 2}

Best analytic distribution, platform B =

{0 10 0 0 0}

This methodology applies to any situation

- I have a homogenous commodity IaaS and want to know whether to buy invest in a novel architecture
- I have tailored IaaS and want to know whether to invest in commodity IaaS
- I have both tailored and commodity IaaS and want to know what future distribution of funds maximizes throughtput
- Any platform, and set of analytics

- Need capital cost, O&M cost, and job runtimes per platform
- Need existing infrastructure, job list, and amount of money to be invested
- Output: which platform acquisition maximizes mission throughput

What I’m looking for

- **Measurements** from current platforms (time-to-solution for every analytic) over some time period
- Candidate platforms for acquisition
- Cost of system, number of racks or nodes, both capital and O&M

BACKUP

Attack the assumptions: space-time trade-off

$$tts = \frac{\sum_i (compute\ unit * time)_i}{number\ of\ compute\ units}$$

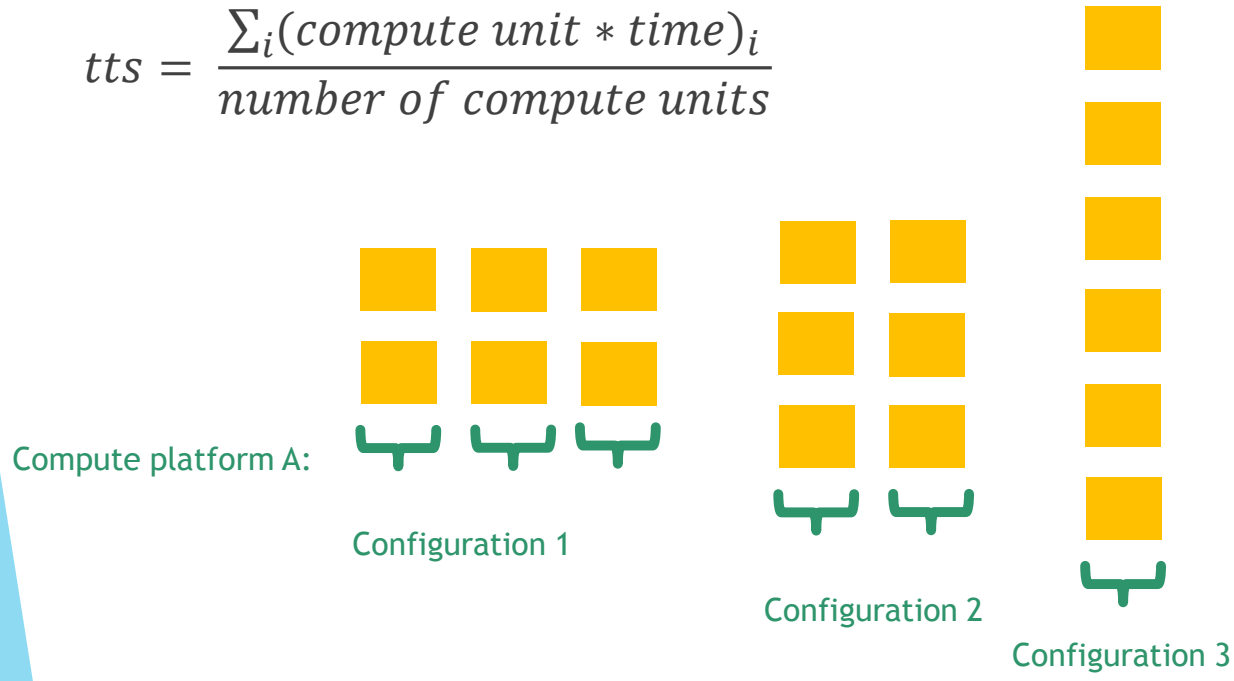

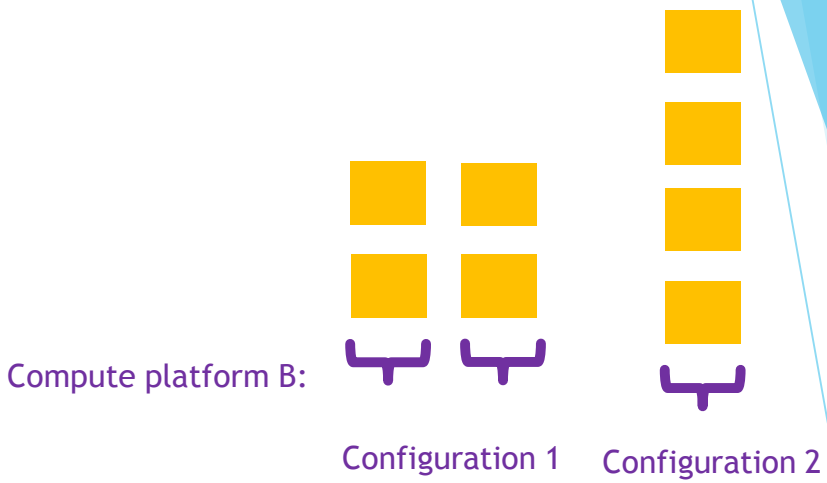

- Space-time trade-off of real applications is not linear; see Amdahl's law
- Scaling depends on how much data is exchanged; synchronous communication

Attack the assumptions: space-time trade-off

- Each job has a scaling curve
- A scheduler searches for the layout of jobs on N compute resources which minimizes time-to-solution
- Scheduler output would yields the time-to-solution per platform

Attack the assumptions: known tasks; known amount of work

- Projection is needed for the fiscal period being modeled
- We have data for what was done in the previous period
- Extrapolate historical data using a scaling factor

Getting faster doesn't matter; the problem is constrained by meatspace

- Investigating different platform architectures wrt time-to-solution

- Providing users more data is not helpful
- Providing users with data faster is not helpful

- Replace "speed-up" with "lower cost to solution"

Stick with homogeneous IaaS, or augment with tailored IaaS?

Goal: balance consumers and resources (avoid both idle and inability to do mission)

Homogeneous environment:

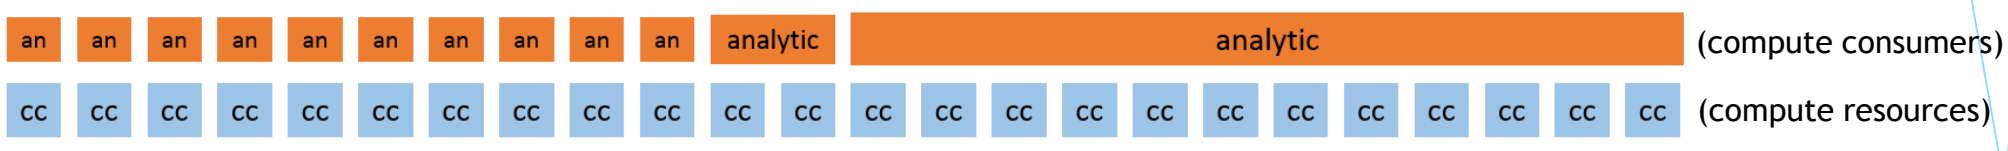

Heterogeneous environment:

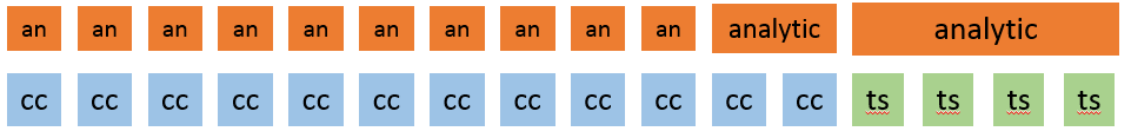

These alternatives could provide equivalent mission throughput;
→ What about $?